




PySpark Optimizations

Spark Shuffles


Map Operation

- Reads one record
- Applies any transformation
- Emits 0..n records
- Each output record depends on exactly one input record

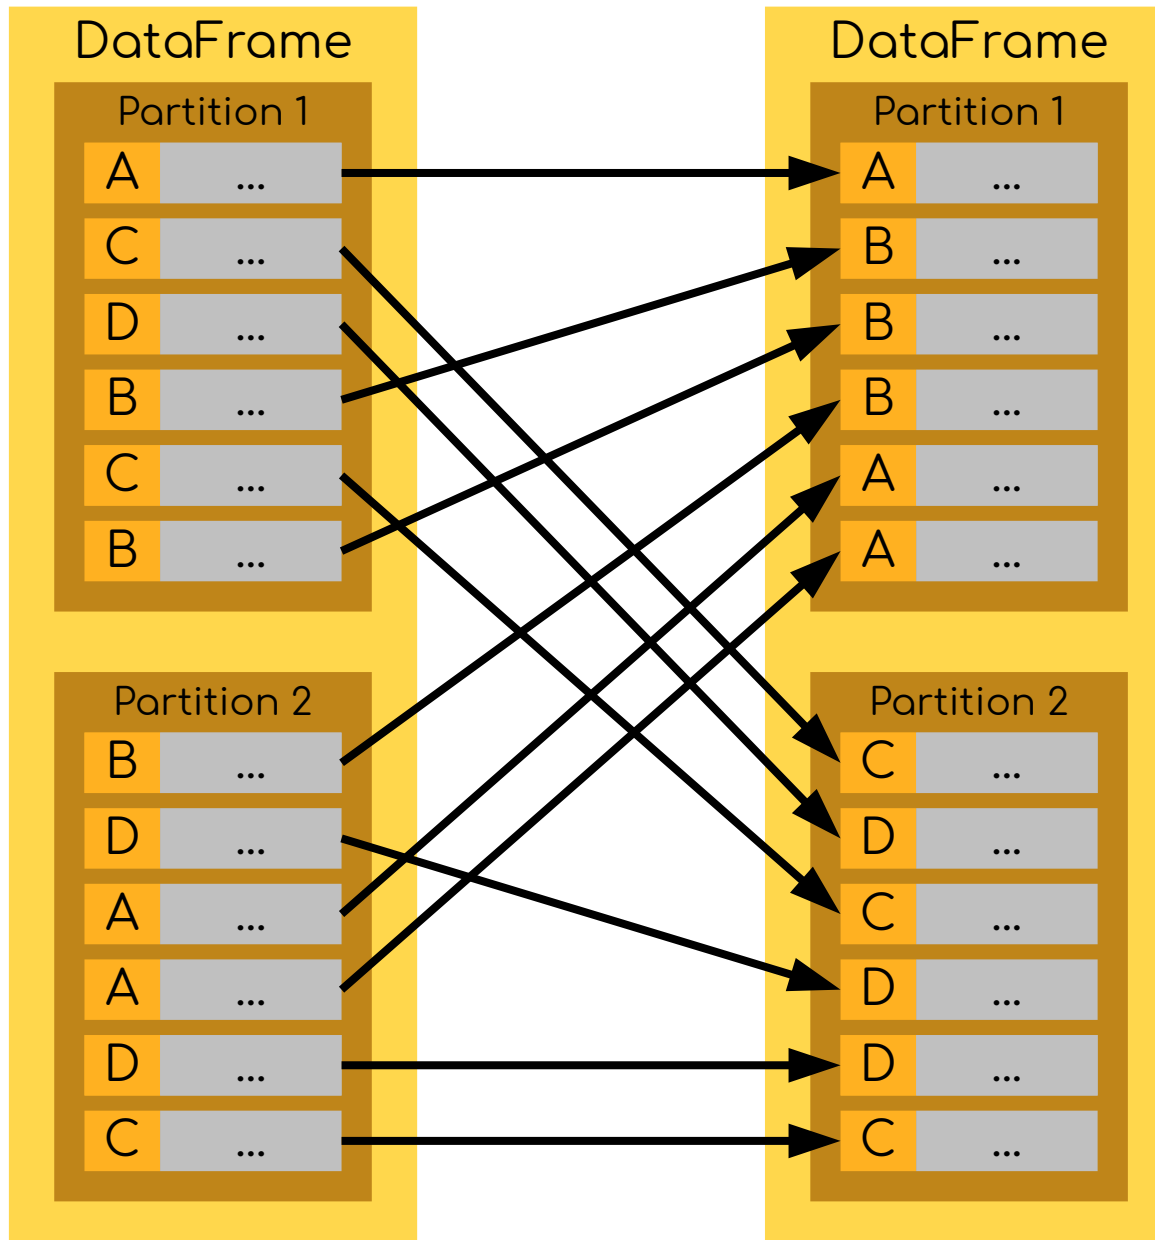


Shuffle Operation


- Collects multiple records with same key
- Results are groups of records



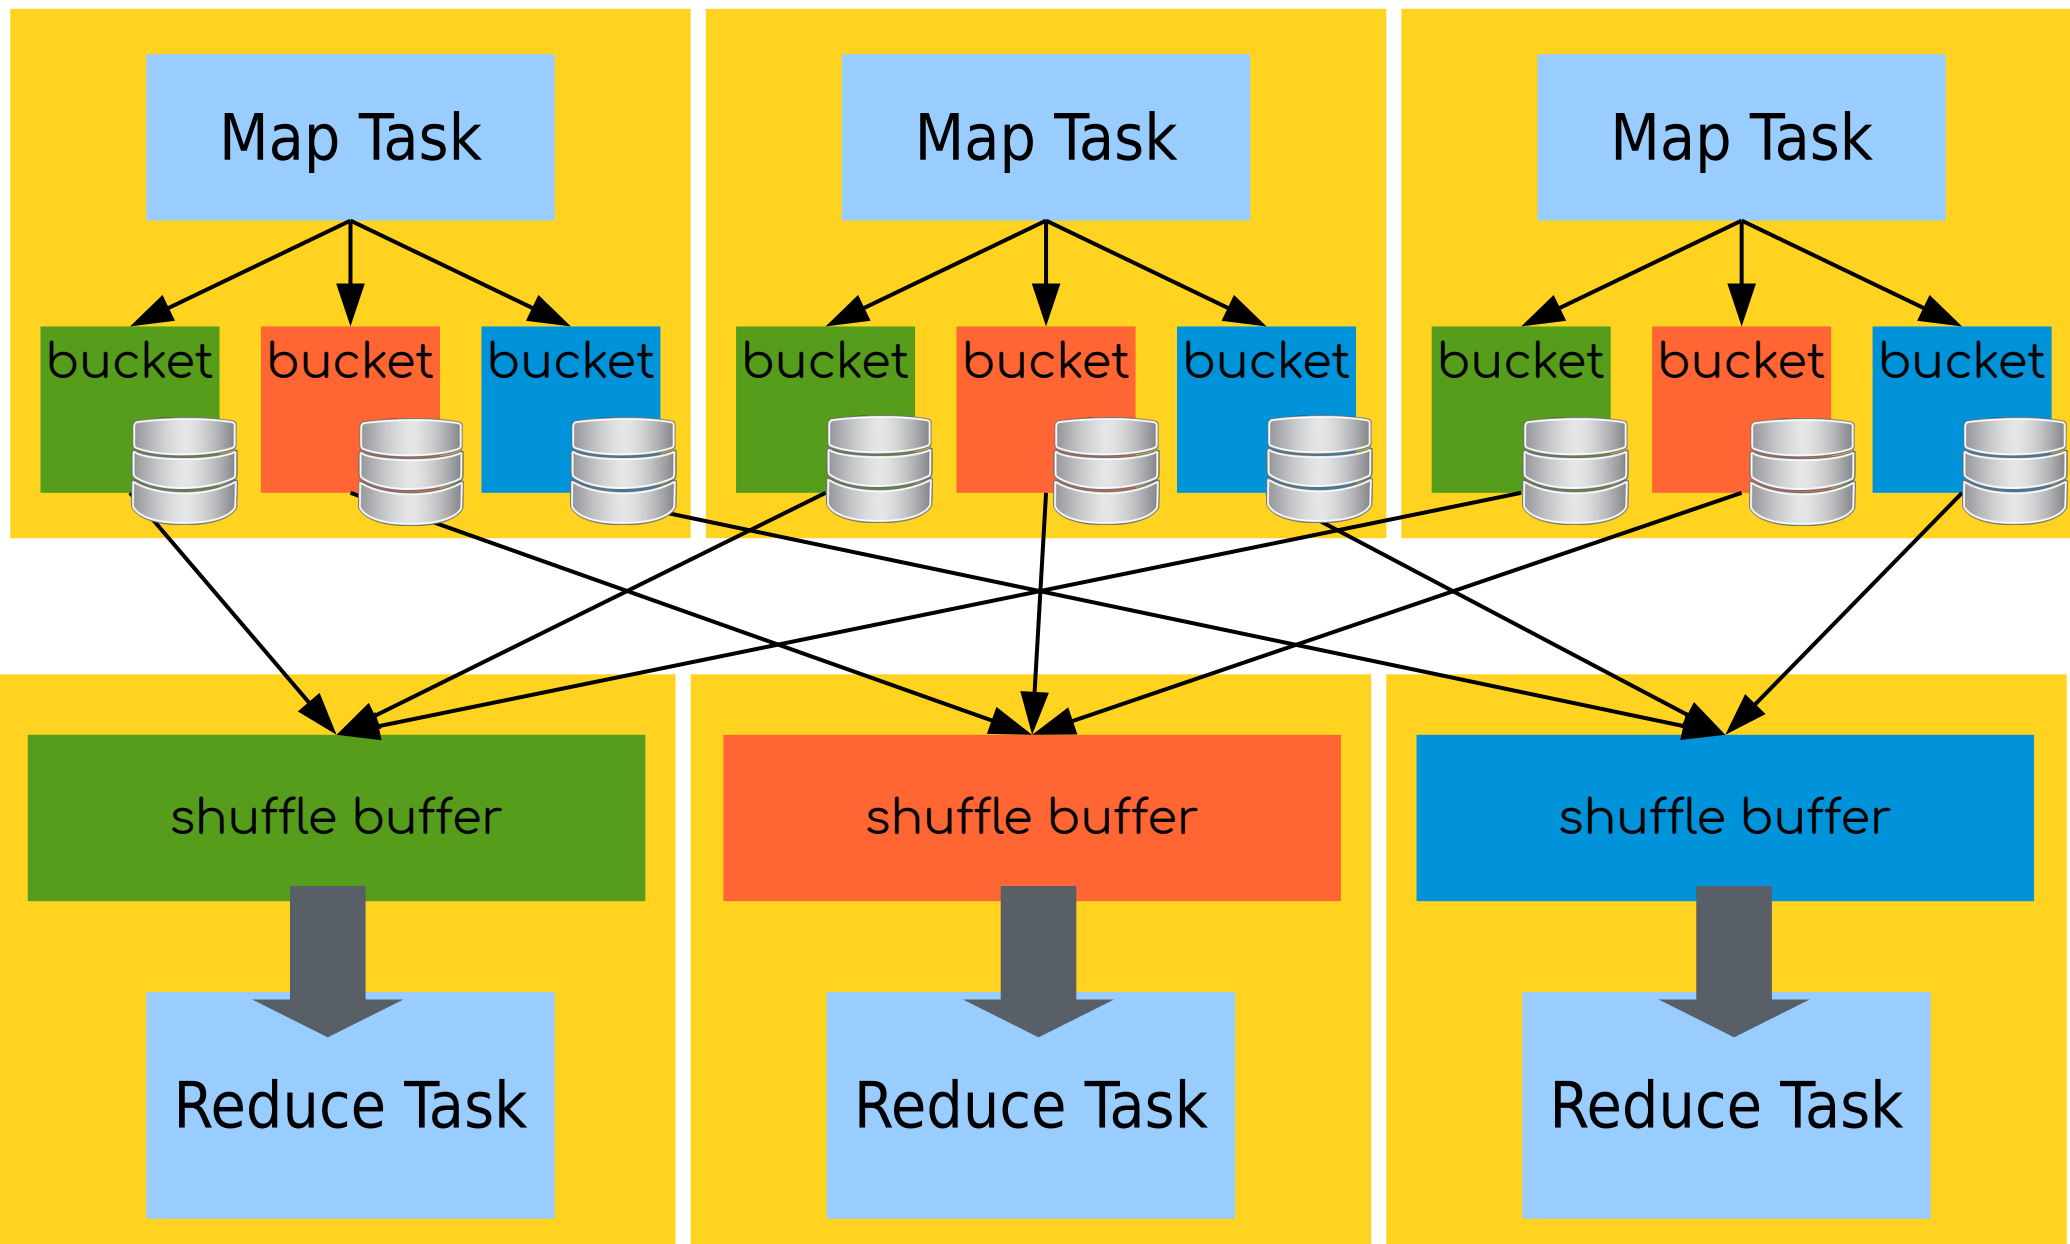
Shuffle Operations




Shuffles in Clusters



Distributed Shuffle

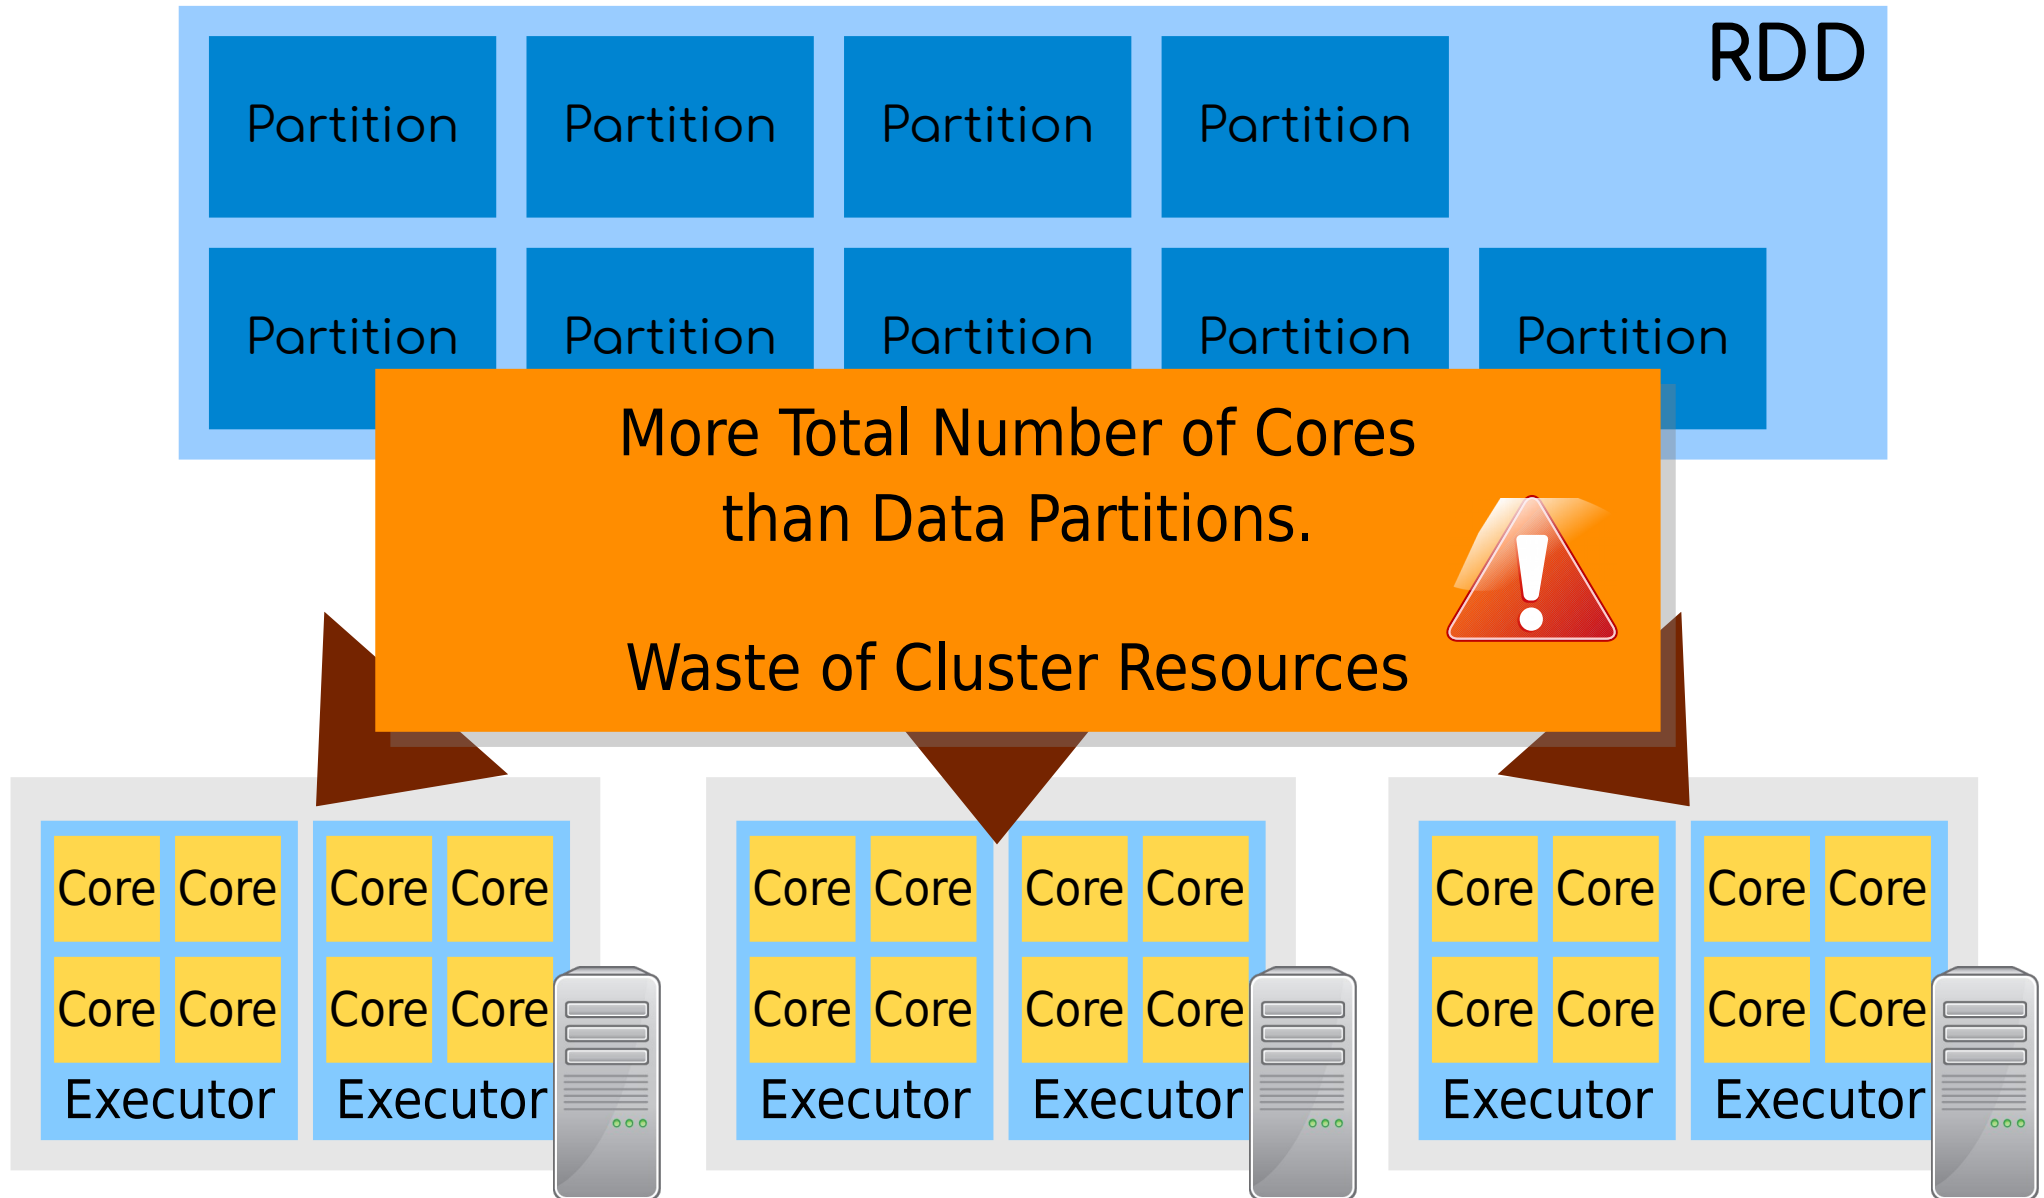


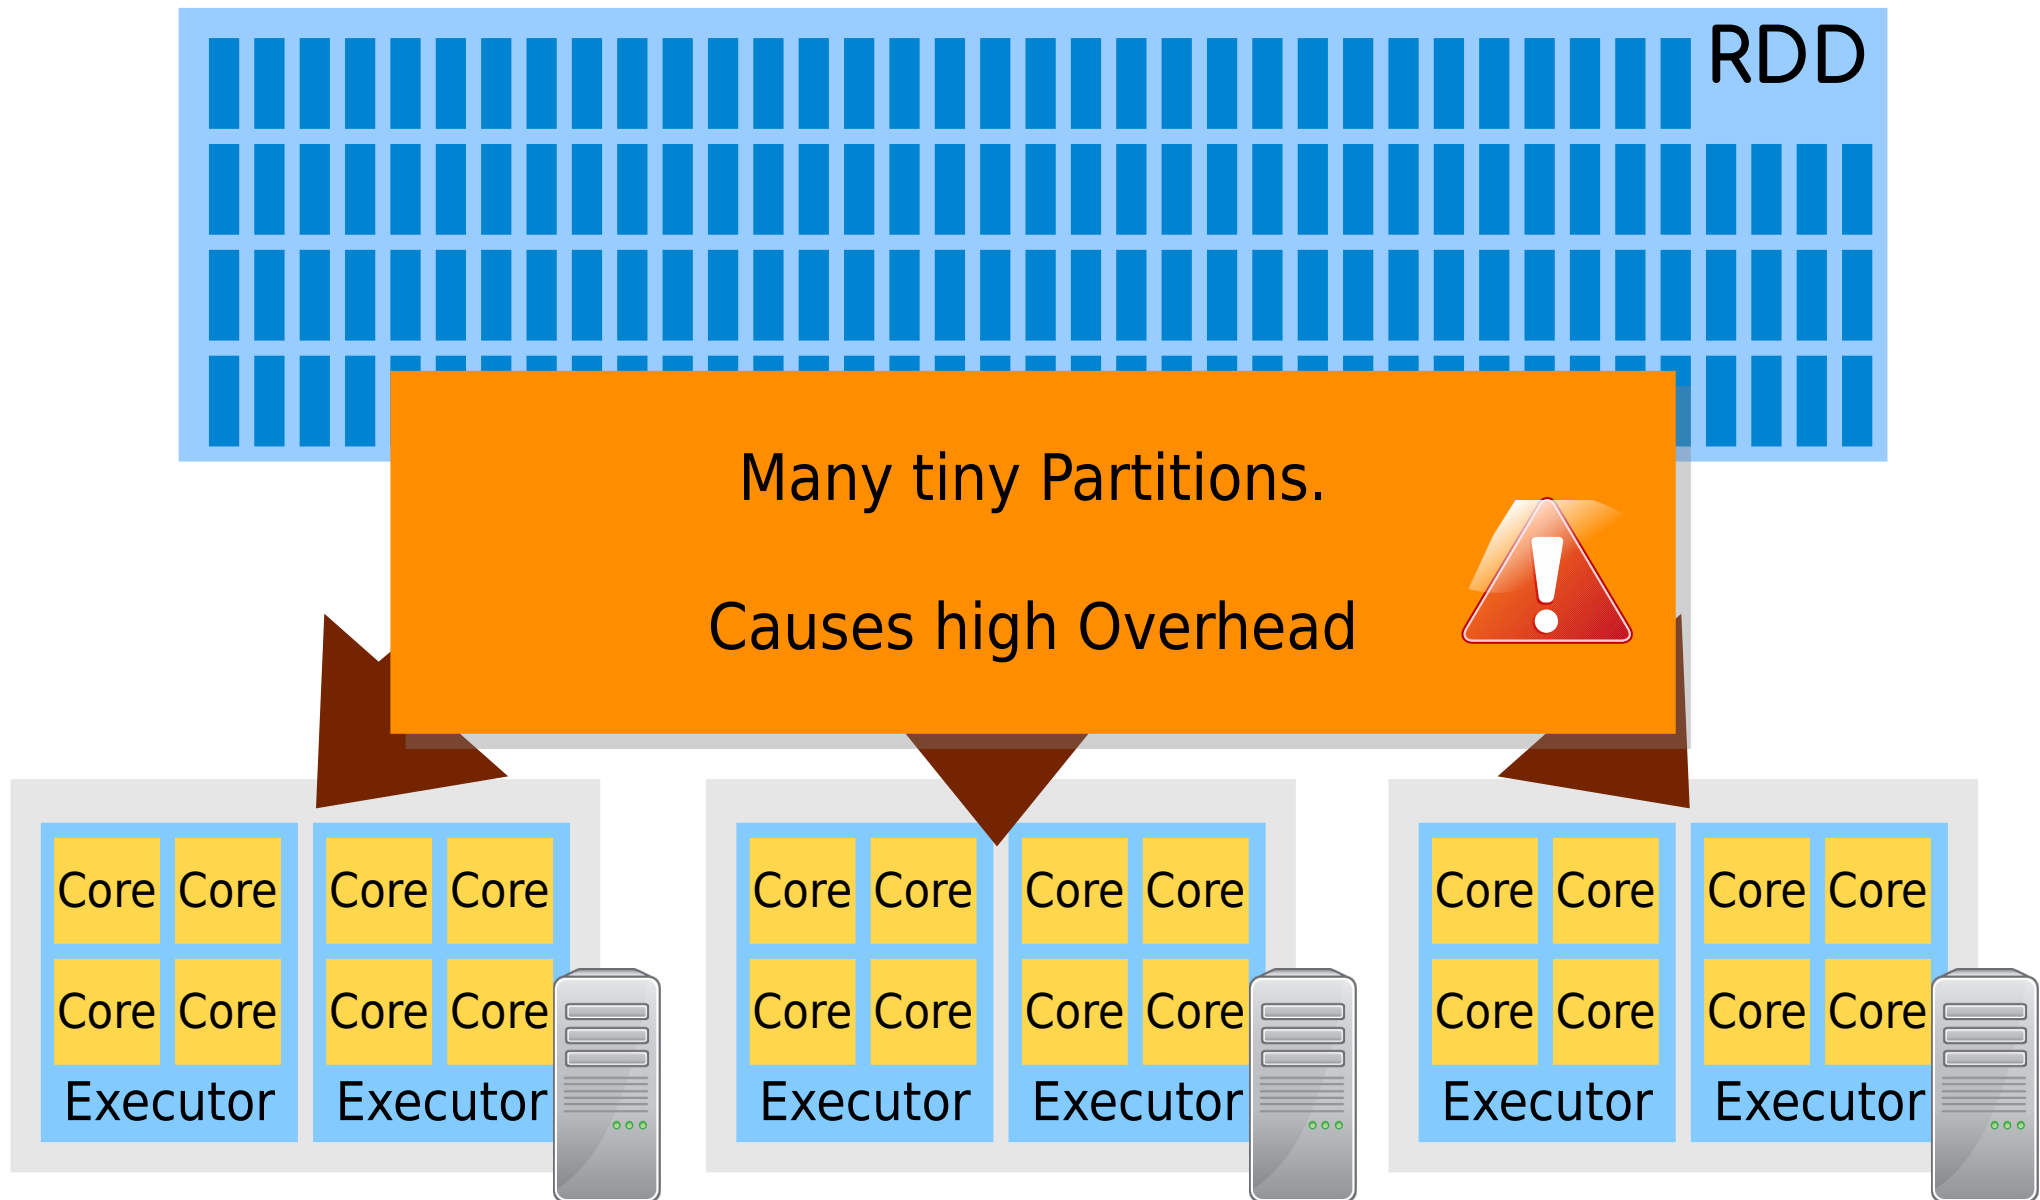
Caching

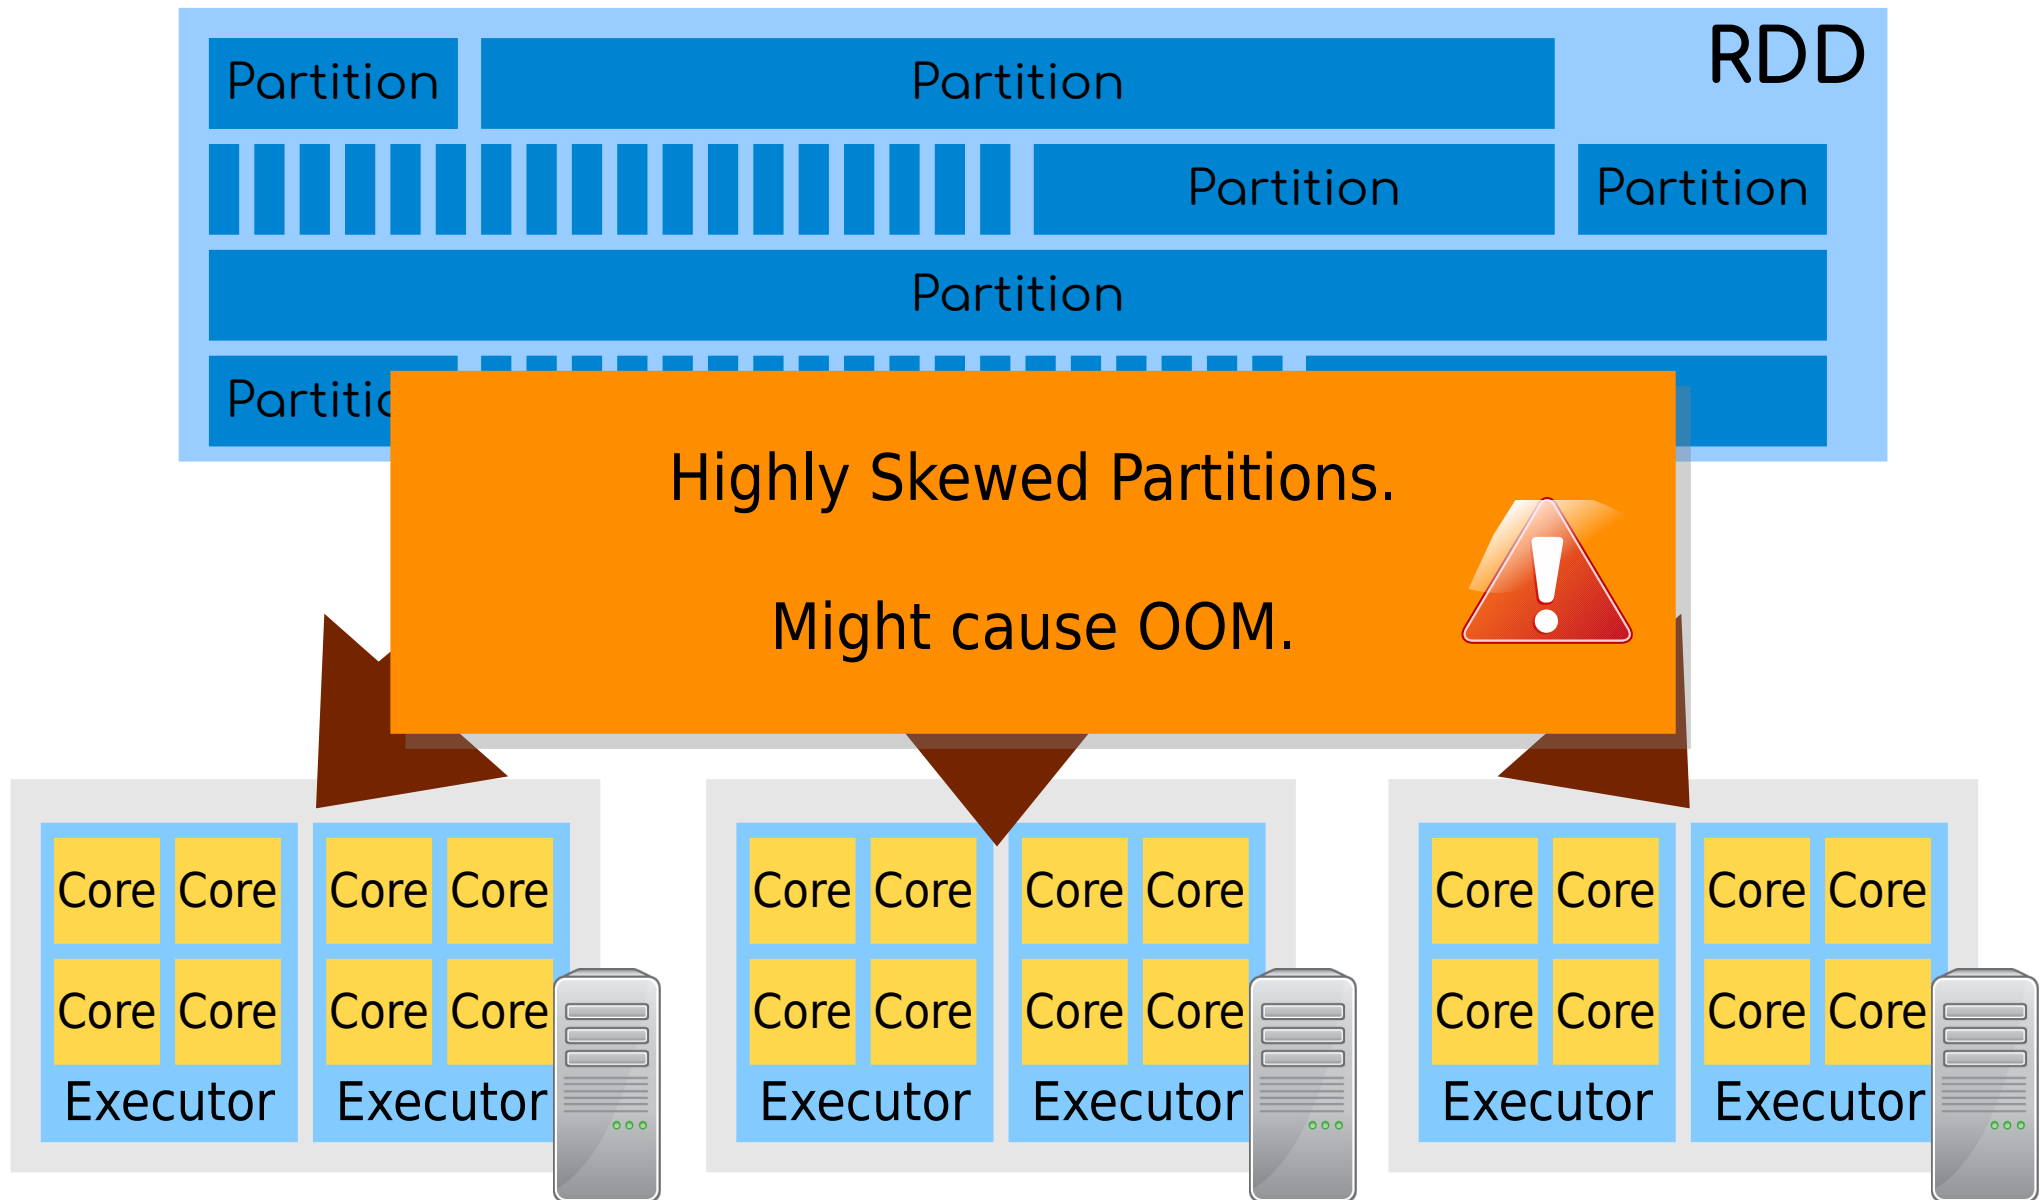


Use `df.cache()` or `df.persist()` for keeping DataFrames that are repeatedly used.

Partitioning








Importance of Good Partitioning

- Large Partitions help to amortize Overhead
- Many Small Partitions help to Utilize Cluster
- But number of Partitions increase overhead!
- Skewed Partitions can cause OOM

No general rule due to many factors:

- Cluster (#Nodes ,RAM, Cores)
- Program (Caching RDDs)
- Data (Skewed Data)




`rdd.repartition(numPartitions)`

Returns a new RDD which has exactly `numPartitions` partitions.

`rdd.coalesce(numPartitions, doShuffle=False)`

Returns a new RDD that is reduced to `numPartitions` partitions.
Returns a narrow dependency, i.e. no shuffle will occur.

`coalesce` is useful for limiting number output files.




`rdd.repartition(numPartitions, *cols)`

Returns a new DataFrame which has exactly `numPartitions` partitions using specified columns for hashing.

`rdd.coalesce(numPartitions)`

Returns a new DataFrame that is reduced to `numPartitions` partitions by logically concatenating partitions.
Returns a narrow dependency, i.e. no shuffle will occur.

`coalesce` is useful for limiting number output files.



`spark.default.parallelism=2`

Default number of partitions of a manually created RDD

`spark.sql.shuffle.partitions=200`

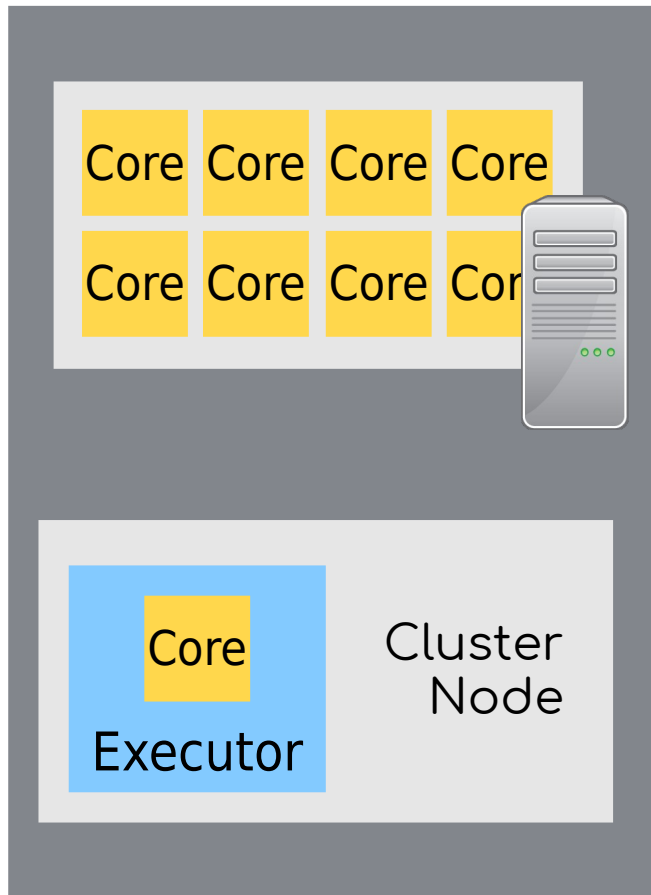
Number of partitions returned by a DataFrame shuffle

`spark.sql.shuffle.partitions`
also controls number output files of SQL operations



Executors

Executor Configuration

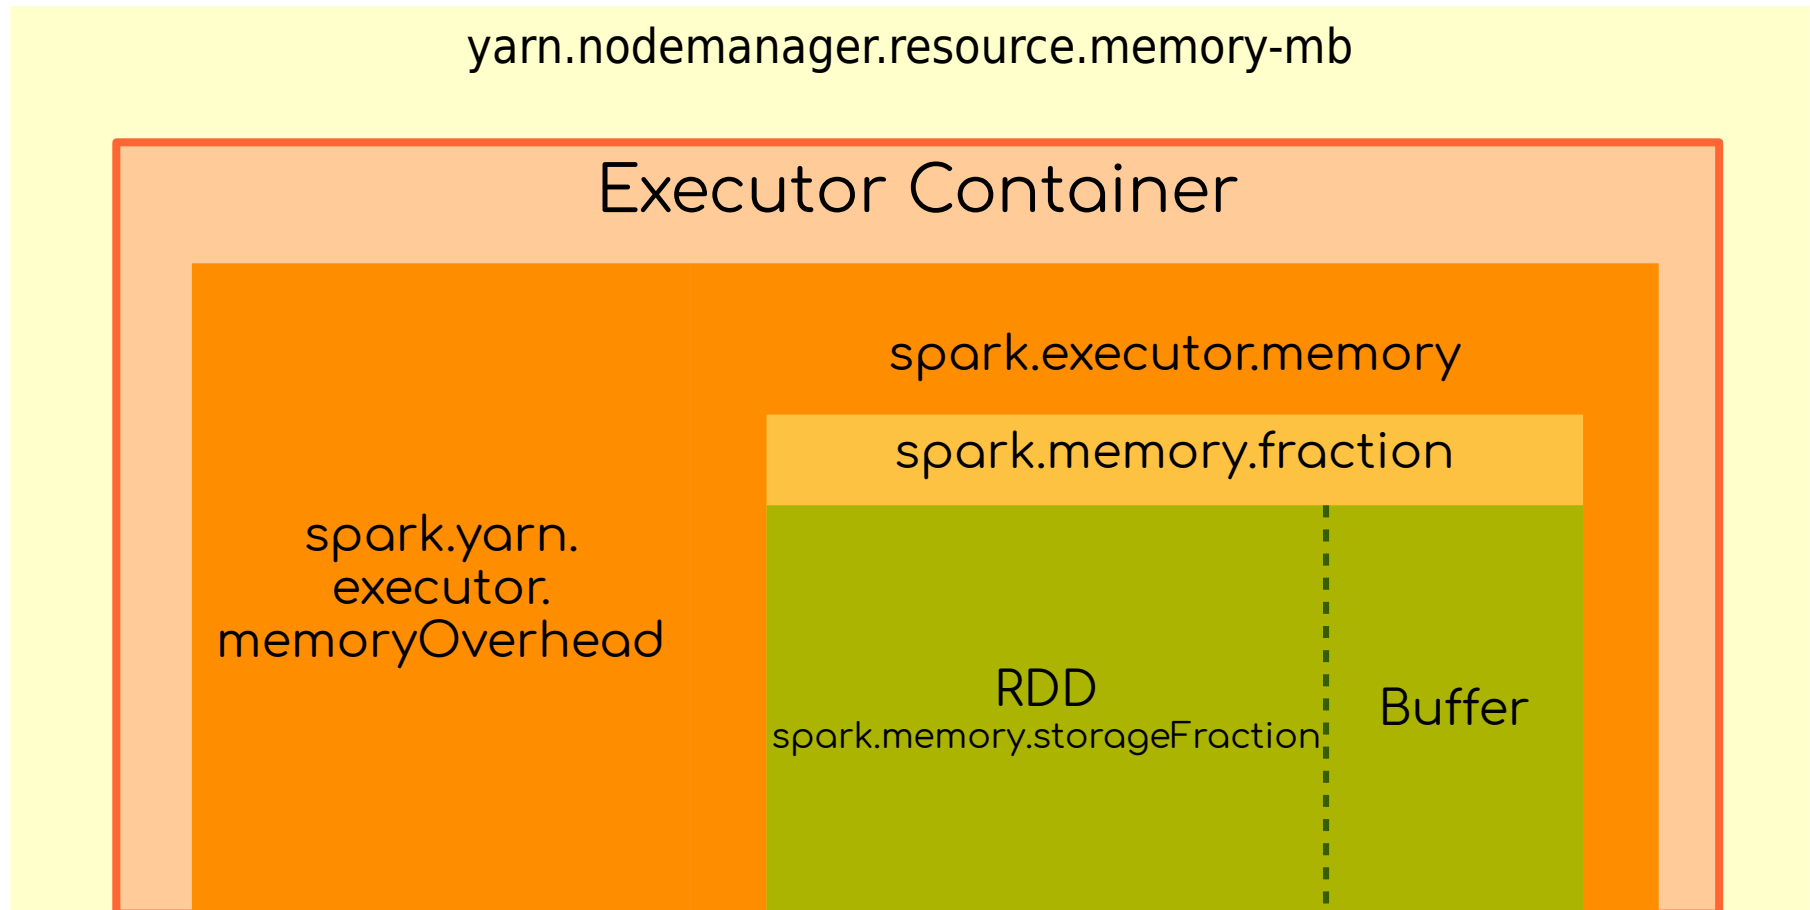


```
spark-submit \  
  <spark options> \  
  <my_program.py> \  
  <my_program_parameters>
```

Command Line Option	Description
<code>--master <master></code>	Master for distributing work in the cluster. Normally „yarn“
<code>--deploy-mode <mode></code>	Deploy driver application into cluster or run it on the client. Normally „client“
<code>--driver-memory <size></code>	Amount of memory for driver program. I.e. 1g, 4g, 256m
<code>--executor-memory <size></code>	Amount of memory per executor. I.e. 1g, 4g, 32g
<code>--executor-cores <number></code>	Number of cores per executor process.
<code>--conf <variable>=<value></code>	Additional configuration properties.

Configuration Option	Default	Description
<code>spark.executor.instances</code>	2	The number of executors.
<code>spark.executor.cores</code>	1	Number of Cores per Executor Process
<code>spark.driver.maxResultSize</code>	1g	
<code>spark.executor.memory</code>		Amount of memory for driver program. I.e. 1g, 4g, 256m
<code>spark.python.worker.memory</code>	512m	Amount of memory to use per python worker process
<code>spark.python.worker.reuse</code>	true	Reuse Python worker or not.

Memory Management



Configuration Option	Default	Description
<code>spark.driver.maxResultSize</code>	1g	
<code>spark.executor.memory</code>		Amount of memory for driver program. I.e. 1g, 4g, 256m
<code>spark.memory.fraction</code>	0.6	Fraction of (heap space - 300MB) used for execution and storage.
<code>spark.memory.storageFraction</code>	0.5	Storage space immune against eviction by execution
<code>spark.yarn.memory.overhead</code>		The amount of off-heap memory (in megabytes) to be allocated per executor.
<code>spark.python.worker.memory</code>	512m	Amount of memory to use per python worker process
<code>spark.python.worker.reuse</code>	true	Reuse Python worker or not.